

# Implement and improve an efficient, layered tape with prefetching capabilities

...

Aditi Milind Joshi

# The Tape

In `clad`, the tape is a stack-like data structure that stores intermediate values in reverse mode AD during the forward pass for use during the backward (gradient) pass.

---

# Old Implementation

Clad used a monolithic memory buffer, which, when full while pushing, reallocated double the capacity.

```
/// Move values from old to new storage.
CUDA_HOST_DEVICE T* AllocateRawStorage(std::size_t _capacity) {
#ifndef CUDAACC
    // Allocate raw storage (without calling constructors of T) of new capacity.
    T* new_data = static_cast<T*>(::operator new(_capacity * sizeof(T)));
#else
    T *new_data =
        static_cast<T*>(::operator new(_capacity * sizeof(T), std::nothrow));
#endif
    return new_data;
}

/// Add new value of type T constructed from args to the end of the tape.
template <typename... ArgsT>
CUDA_HOST_DEVICE void emplace_back(ArgsT&&... args) {
    if (_size >= _capacity)
        grow();
    ::new (const_cast<void*>(static_cast<const volatile void*>(end())))
        T(std::forward<ArgsT>(args)...);
    _size += 1;
}
```

```
private:
    // Copies the data from a storage to another.
    // Implementation taken from std::uninitialized_copy
    template <class InputIt, class NoThrowForwardIt>
    CUDA_HOST_DEVICE void MoveData(InputIt first, InputIt last,
                                   NoThrowForwardIt d_first) {
        NoThrowForwardIt current = d_first;
        // We specifically add and remove the CV qualifications here so that
        // cases where NoThrowForwardIt is CV qualified, we can still do the
        // allocation properly.
        for (; first != last; ++first, ++current) {
            ::new (const_cast<void*>(
                static_cast<const volatile void*>(clad_addressof(*current))))
                T(std::move(*first));
        }
    }

    // Initial capacity (allocated whenever a value is pushed into empty tape).
    constexpr static std::size_t _init_capacity = 32;
    CUDA_HOST_DEVICE void grow() {
        // If empty, use initial capacity.
        if (!_capacity)
            capacity = _init_capacity;
        else
            // Double the capacity on each reallocation.
            capacity *= 2;
        T* new_data = AllocateRawStorage(_capacity);

        if (!new_data) {
            // clean up the memory mess just in case!
            destroy(begin(), end());
            printf("Allocation failure during tape resize! Aborting.\n");
            trap(EXIT_FAILURE);
        }
    }
```

# Midterm Progress - Slab-based Tape with SBO

- The new tape follows a slab-based structure with small buffer optimization.
- Slab size and buffer size are configurable template parameters.
- Slab size 1024 and buffer size 64 performed the best and were set as default values.

```
/// A dynamic slab-based vector-like container with Small Buffer Optimization
/// (SBO) primarily used for storing values in reverse mode AD. Stores elements
/// in a static buffer first, then falls back to dynamically allocated linked
/// slabs if capacity exceeds SBO_SIZE.
template <typename T, std::size_t SBO_SIZE = 64, std::size_t SLAB_SIZE = 1024>
class tape {
public:
    // A block of contiguous storage allocated dynamically when SBO capacity is
    // exceeded.
    struct Slab {
        // std::aligned_storage_t<sizeof(T), alignof(T)> raw_data[SLAB_SIZE];
        // For now use the implementation below as above implementation is not
        // supported by c++11
        alignas(T) char raw_data[SLAB_SIZE * sizeof(T)];
        Slab* next;
        CUDA_HOST_DEVICE Slab() : next(nullptr) {}
        CUDA_HOST_DEVICE T* elements() {
            #if __cplusplus >= 201703L
                return std::launder(reinterpret_cast<T*>(raw_data));
            #else
                return reinterpret_cast<T*>(raw_data);
            #endif
        }
        // std::aligned_storage_t<sizeof(T), alignof(T)> in static buffer[SBO_SIZE];
        // For now use the implementation below as above implementation is not
        // supported by c++11
        alignas(T) char m_static_buffer[SBO_SIZE * sizeof(T)];
        bool m_using_sbo = true;
        Slab* m_head = nullptr;
        std::size_t m_size = 0;
    };
    CUDA_HOST_DEVICE T* sbo_elements() {
        #if __cplusplus >= 201703L
            return std::launder(reinterpret_cast<T*>(m_static_buffer));
        #else
            return reinterpret_cast<T*>(m_static_buffer);
        #endif
    }
    CUDA_HOST_DEVICE const T* sbo_elements() const {
        #if __cplusplus >= 201703L
            return std::launder(reinterpret_cast<const T*>(m_static_buffer));
        #else
            return reinterpret_cast<const T*>(m_static_buffer);
        #endif
    }
};
```

```
/// Add new value of type T constructed from args to the end of the tape.
template <typename... Args>
CUDA_HOST_DEVICE void emplace_back(Args&&... args) {
    if (m_size < SBO_SIZE) {
        // Store in SBO buffer
        ::new (const_cast<void*>(static_cast<const volatile void*>(
            m_sbo_elements() + m_size))) T(std::forward<ArgsT>(args)...);
    } else {
        // Transition to dynamic storage if needed
        if (m_using_sbo)
            m_using_sbo = false;
        // Allocate new slab if required
        if ((m_size - SBO_SIZE) % SLAB_SIZE == 0) {
            Slab* new_slab = new Slab();
            if (!m_head) {
                m_head = new_slab;
            } else {
                Slab* last = m_head;
                while (last->next)
                    last = last->next;
                last->next = new_slab;
            }
        }
        // Find correct slab for element
        Slab* slab = m_head;
        std::size_t idx = (m_size - SBO_SIZE) / SLAB_SIZE;
        while (idx--)
            slab = slab->next;
        // Construct element in-place
        ::new (const_cast<void*>(static_cast<const volatile void*>(
            slab->elements() + ((m_size - SBO_SIZE) % SLAB_SIZE))) T(std::forward<ArgsT>(args)...));
        m_size++;
    }
}
```

# Fixes in Benchmarks & Tape Improvements

## Benchmark Fixes

- Removed Iterations(1) to get better estimate of the benchmarks.
- Fixed Memory Manager Counters.
- Added DoNotOptimize() to prevent compiler from optimizing out the pop function.

## Tape Improvements

- Added Tail Pointer reducing the push and pop function runs from  $O(n)$  to  $O(1)$ .
- Changed tape to doubly linked list.
- Reused slabs which were previously allocated.

# Updated Code

```
/// A dynamic slab-based vector-like container with Small Buffer Optimization
/// (SBO), primarily used for storing values in reverse-mode AD. Stores elements
/// in a static buffer first, then falls back to dynamically allocated linked
/// slabs if capacity exceeds SBO.
template <typename T, std::size_t SBO_SIZE = 64, std::size_t SLAB_SIZE = 1024,
          bool is_multithread = false>
class tape_impl {
    /// A block of contiguous storage allocated dynamically when SBO capacity is
    /// exceeded.
    struct Slab {
        // std::aligned_storage_t<sizeof(T), alignof(T)> raw_data[SLAB_SIZE];
        // For now use the implementation below as above implementation is not
        // supported by c++11
        alignas(T) char raw_data[SLAB_SIZE * sizeof(T)];
        Slab* prev;
        Slab* next;
        CUDA_HOST_DEVICE Slab() : prev(nullptr), next(nullptr) {}
        CUDA_HOST_DEVICE T* elements() {
            #if __cplusplus >= 201703L
                return std::launder(reinterpret_cast<T*>(raw_data));
            #else
                return reinterpret_cast<T*>(raw_data);
            #endif
        }
    };
    // std::aligned_storage_t<sizeof(T), alignof(T)> m_static_buffer[SBO_SIZE];
    // For now use the implementation below as above implementation is not
    // supported by C++11
    alignas(T) char m_static_buffer[SBO_SIZE * sizeof(T)];

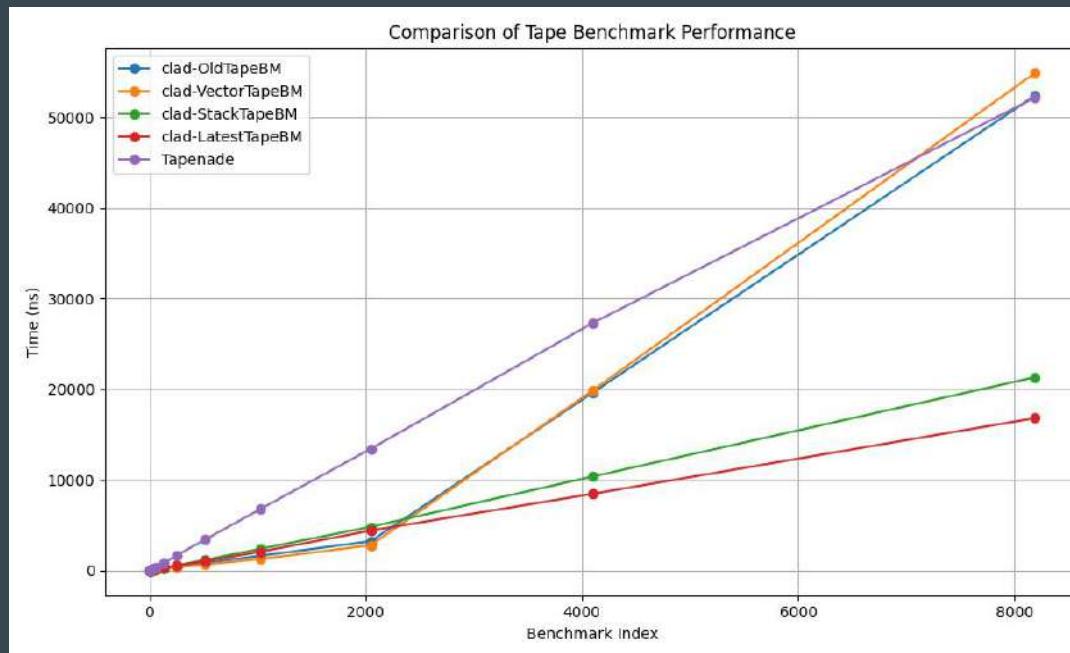
    Slab* m_head = nullptr;
    Slab* m_tail = nullptr;
    std::size_t m_size = 0;
    std::size_t m_capacity = SBO_SIZE;
```

```
/// Add new value of type T constructed from args to the end of the tape.
template <typename... Args>
CUDA_HOST_DEVICE void emplace_back(ArgsT&... args) {
    if (m_size < SBO_SIZE) {
        // Store in SBO buffer
        ::new (const cast<void*>(static cast<const volatile void*>(
            sbo_elements() + m_size))) T(std::forward<ArgsT>(args)...);
    } else {
        const auto offset = (m_size - SBO_SIZE) % SLAB_SIZE;
        // Allocate new slab if required
        if (!offset) {
            if (m_size == m_capacity) {
                Slab* new_slab = new Slab();
                if (!m_head)
                    m_head = new_slab;
                else {
                    m_tail->next = new_slab;
                    new_slab->prev = m_tail;
                }
                m_capacity += SLAB_SIZE;
            }
            if (m_size == SBO_SIZE)
                m_tail = m_head;
            else
                m_tail = m_tail->next;
        }
        // Construct element in-place
        ::new (const cast<void*>(static cast<const volatile void*>(
            m_tail->elements() + offset))) T(std::forward<ArgsT>(args)...);
    }
    m_size++;
}
```

```
/// Remove the last value from the tape.
CUDA_HOST_DEVICE void pop_back() {
    assert(m_size);
    m_size--;
    if (m_size < SBO_SIZE)
        destroy_element(sbo_elements() + m_size);
    else {
        std::size_t offset = (m_size - SBO_SIZE) % SLAB_SIZE;
        destroy_element(m_tail->elements() + offset);
        if (offset == 0) {
            if (m_tail != m_head)
                m_tail = m_tail->prev;
        }
    }
}
```

# New Benchmarks

- Tested against the old tape, std::vector, std::stack and tapenade.



# Thread Safety

- Added thread-safe tape access functions with mutex locking mechanism to allow concurrent access.

```
/// Thread safe tape access functions with mutex locking mechanism
#ifndef __CUDAACC__
// Add value to the end of the tape, return the same value.
template <typename T, std::size_t SBO_SIZE = 64, std::size_t SLAB_SIZE = 1024,
          typename... ArgsT>
T push(tape<T, SBO_SIZE, SLAB_SIZE, /*is_multithreaded*/true>& to,
       ArgsT... val) {
    std::lock_guard<std::mutex> lock(to.mutex());
    to.emplace_back(std::forward<ArgsT>(val)...);
    return to.back();
}

// A specialization for C arrays
template <typename T, typename U, std::size_t N, std::size_t SBO_SIZE = 64,
          std::size_t SLAB_SIZE = 1024>
void push(tape<T[N], SBO_SIZE, SLAB_SIZE, /*is_multithreaded*/true>& to,
          const U& val) {
    std::lock_guard<std::mutex> lock(to.mutex());
    to.emplace_back();
    std::copy(std::begin(val), std::end(val), std::begin(to.back()));
}

// Remove the last value from the tape, return it.
template <typename T, std::size_t SBO_SIZE = 64, std::size_t SLAB_SIZE = 1024>
T pop(tape<T, SBO_SIZE, SLAB_SIZE, /*is_multithreaded*/true>& to) {
    std::lock_guard<std::mutex> lock(to.mutex());
    T val = std::move(to.back());
    to.pop_back();
    return val;
}

// A specialization for C arrays
template <typename T, std::size_t N, std::size_t SBO_SIZE = 64,
          std::size_t SLAB_SIZE = 1024>
void pop(tape<T[N], SBO_SIZE, SLAB_SIZE, /*is_multithreaded*/true>& to) {
    std::lock_guard<std::mutex> lock(to.mutex());
    to.pop_back();
}

// Access return the last value in the tape.
template <typename T, std::size_t SBO_SIZE = 64, std::size_t SLAB_SIZE = 1024>
T& back(tape<T, SBO_SIZE, SLAB_SIZE, /*is_multithreaded*/true>& of) {
    std::lock_guard<std::mutex> lock(of.mutex());
    return of.back();
}
#endif
```

# Ongoing Progress - Multilayer Tape

## Multilayer Tape Storage

- Offloading mechanism to offload slabs to disk and load slabs from disk to memory.
- Keep only the last N slabs in memory at a time and offload the rest.
- One slab space where slabs are loaded if element to be loaded is not in memory.

# Progress

## Current Progress

- Benchmark script to compare two revisions (PR #1394)
  - Modified tape structure to slab-based (PR #1404)
  - Added small buffer optimization (PR #1404)
  - Enhanced benchmarks (PR #1404 and PR #1533)
  - Made improvements to tape (PR #1531 and PR #1633)
  - Added thread safety (PR #1470)
-

# Thank You!